

# **An Implementation of User-Level Processes using Address Space Sharing**

**IPDPS/RADR 2020 (virtual meeting)**

**May 18, 2020**

**Atsushi Hori, Balazs Gerofi ,Yutaka Ishikawa**

**Riken CCS (R-CCS)**

**JAPAN**

# Outline

## 1. Put an end to the long-term discussion

- Kernel-Level Thread (KLT) vs. User-Level Thread (ULT)
  - Advantages and disadvantages

## 2. Challenges

### 1. Bi-Level Thread

- To take the best of the two worlds
- A ULT can be a KLT and vice versa
- User-level (thread) context switching
- Blocking system-call can be called as a kernel-level thread

### 2. User-Level Process


- Combining Bi-Level Thread with [Address Space Sharing](#)
  - Process context switching at user-level

## 3. Evaluation


# Bi-Level Thread

# Re-thinking Thread Models


- Thread models (1:1, N:1, and M:N)
  - KC: Kernel Context, UC: User Context



(a) Kernel-Level Threads (1:1)



(b) User-Level Threads (N:1)



(c) User-Level Threads (M:N)

- What if KCs and UCs in 1:1 model can be decoupled and coupled again ?
  - The 1:1 model and M:N ( $M=N$ ) model can be interchangeable

# The Idea of decoupling



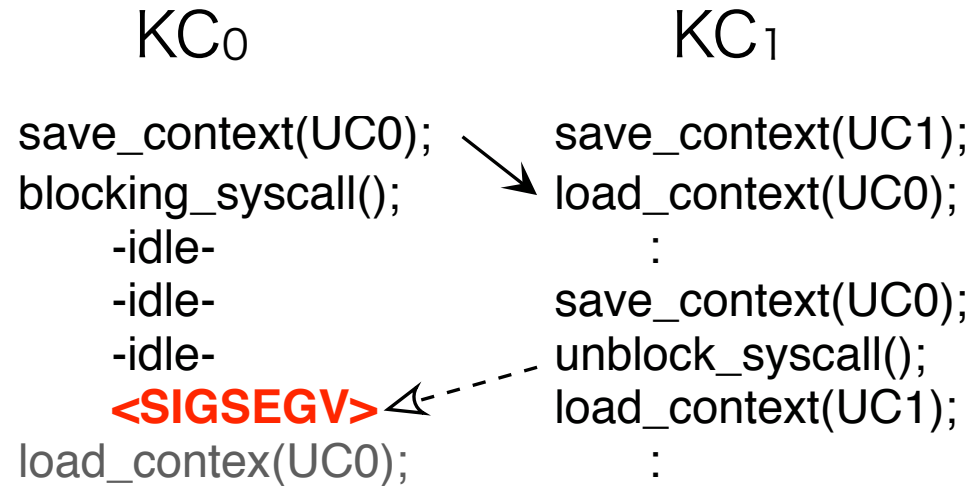


Fig. 2. Coupling and decoupling of UCs and KCs

# Decoupling and Coupling

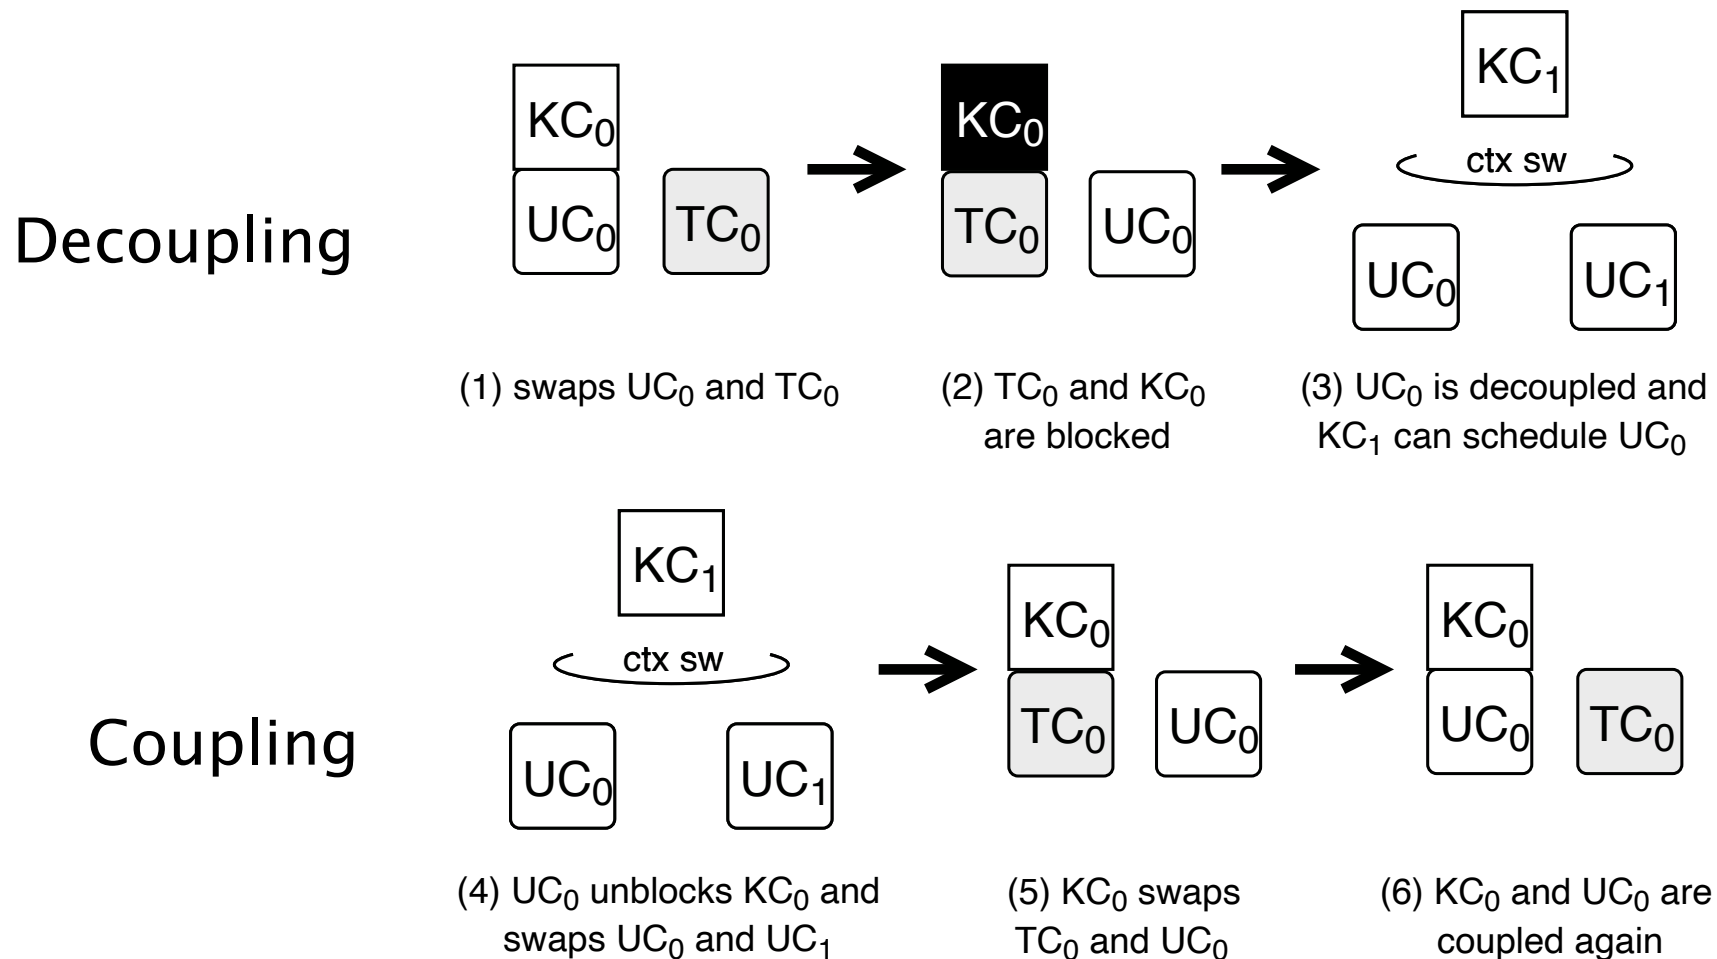
- What if a UC is decoupled from KC ?
  - Decoupled UC can be scheduled by another KC
  - What happens on the decoupled KC ?
    - It has nothing to do (idling or blocked in some way)
  - This is transition for a KLT (Kernel-Level Thread) to be a ULT (User-Level Thread)
- What if the decoupled UC wants to be coupled again ?
  - The idling KC now schedules the UC
  - This is the transition for a ULT to be a KLT
- However, KC must always be associated with a UC
  - A KC cannot be idling without a UC
  - But the UC has to be decoupled so that it can be scheduled by another KC...

# Issue in Decoupling and Coupling



# Trampoline Context

- This problem can be resolved by introducing another small context ([Trampoline Context, TC](#))





# Resolving Blocking System-call Issue

- Issue

- When a ULT calls a blocking system-call,
  - the scheduling KC is also blocked, and
  - the other eligible-to-run ULTs have no chance to be scheduled

- Solution by using coupling and decoupling

- Assumption:
  - A ULT is going to call a blocking system-call
  - The ULT was firstly created with KCo
  - The ULT is already decoupled and scheduled by the KCs
- 1. before the ULT calls the system-call, it is coupled with KCo
- 2. (KCs can schedule the other ULTs)
- 3. the ULT becomes a KLT and it calls the system-call
- 4. after returning from the system-call, it is decoupled so that KCs can schedule it (becoming ULT again)

# User-Level Process

# Address Space Sharing

- What is Address Space Sharing (ASS) technique ?
  - “Processes” share the same address space
  - Here “process” is defined as an execution entity having privatized static variables, and ASS “processes” may be derived from different programs
    - threads share all static variables, and threads are derived from the same program
  - ASS is different from POSIX shared memory (PSM)
    - ASS share the whole address space (and a page table)
    - PSM shares only some specific memory pages
- Process-in-Process (PiP)
  - Pure user-level implementation of ASS



A. Hori, M. Si, B. Geroft, M. Takagi, J. Dayal, P. Balaji, and Y. Ishikawa, “Process-in-process: Techniques for practical address-space sharing,” in *Proceedings of the 27th International Symposium on High-Performance Parallel and Distributed Computing*, ser. HPDC 18.

# BLT + ASS = User-Level Process (ULP)

- ASS allows for a process to context-switch one to the other at user-level => Fast context switching
- The difference between ULT and ULP
  - Threads share most OS kernel resources while processes do not
  - Example: getpid()
    - threads have the same PID
    - each process has its own unique PID
- In ULP
  - System-call consistency can be preserved
    - by using the same (de)coupling technique in BLT
  - Thread Local Storage (TLS) must also be switched when switching contexts
    - In most ULT implementations, TLS switching is ignored

# Evaluation

# Evaluation Results

- Machines: Wallaby – x86\_64, Albireo – AArch64

TABLE III  
CONTEXT SWITCH AND LOAD TLS

	Wallaby		Albireo
	Time [Sec]	Cycles	Time [Sec]
Context Sw.	3.34E-8	86	2.45E-8
Load TLS	1.09E-7	284	2.50E-9

On x86\_64 CPUs a system-call is required to switch TLS

TABLE V  
TIME OF `getpid()`

	Wallaby		Albireo
	Time [Sec]	Cycles	Time [Sec]
Linux	6.71E-8	174	3.85E-7
ULP-PiP: BUSYWAIT	1.33E-6	3452	2.71E-6
ULP-PiP: BLOCKING	2.91E-6	6172	4.48E-6

BUSYWAIT/BLOCKING:  
Idling ways of KC w/ TC

In ULP-PiP cases, `getpid()` call is wrapped by `pip_couple()` and `pip_decouple()`




Fig. 8. Comparison of Overlap Ratios

ULP-PiP:  
`pip_couple();`  
`open();`  
`write(Size);`  
`close();`  
`pip_decouple();`

AIO-return:  
`open();`  
`aio_write(Size);`  
`do {`  
    `aio_return();`  
`} while();`  
`close();`

AIO-suspend:  
`open();`  
`aio_write(Size);`  
`aio_suspend();`  
`close();`

# Summary

ULP-PiP will be available at  
<https://github.com/RIKEN-SysSoft>

- Proposing
  - Bi-Level Thread (BLT)
    - Decoupling and coupling UC and KC
    - Trampoline Context to block decoupled KC
    - Able to handle blocking system-calls effectively
  - User-Level Process (ULP) by using Address Space Sharing
    - Switching Thread Local Storage (TLS)
    - System-call consistency
  - Coupling and decoupling can be applied to
    - resolve the blocking system-call issue, and
    - preserve system-call consistency in ULP
- Evaluation (ULP vs. AIO)
  - Coupling and decoupling scheme of ULP-PiP outperforms AIO